

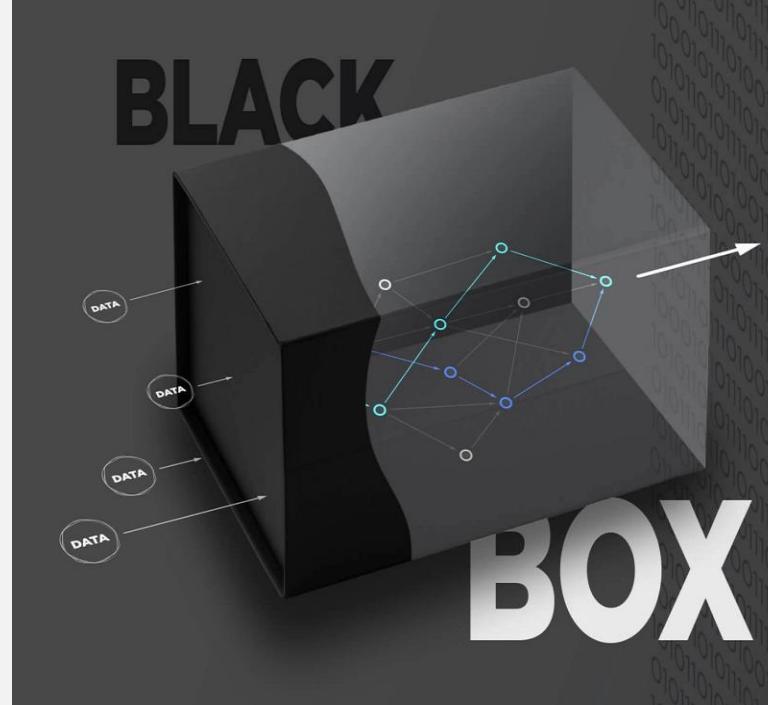
UNBOXING-BLACKBOX

BLACK

BOX

THE LAZY ARTIST-CV

Mātaram Pitrum Vande|
Āchārya Devo Bhava



By P. Srivatsav Reddy
2025121016

TASK0- The Biased Canvas:

EASY TRAIN

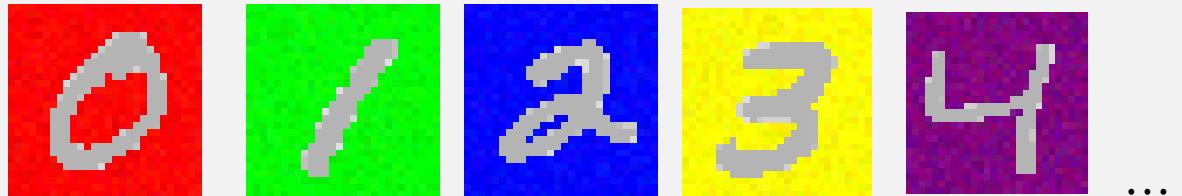
Digit	Dominant Color Count	Total	Percent
0	5615	5923	94.80%
1	6409	6742	95.06%
2	5692	5958	95.54%
3	5835	6131	95.17%
4	5521	5842	94.51%
5	5138	5421	94.78%
6	5642	5918	95.34%
7	5959	6265	95.12%
8	5565	5851	95.11%
9	5657	5949	95.09%

HARD TEST

Digit	Dominant Color Count	Total	Percent
0	0	980	0.00%
1	0	1135	0.00%
2	0	1032	0.00%
3	0	1010	0.00%
4	0	982	0.00%
5	0	892	0.00%
6	0	958	0.00%
7	0	1028	0.00%
8	0	974	0.00%
9	0	1009	0.00%

I made background textured and colored it and foreground strokes gray to prob insane !

Train set: 95 percent each-digit:biased towards same colour



Test set: negation of colors

[Created kaagle datacard :](#)

<https://www.kaggle.com/datasets/poreddysr/c2ty2bcouldspurious>

TASK I - The Cheater:

Before cheating , I used proper STRATIFIED Split in the dataset and used 25percent of data
MNIST Training ALL models out of 60k training 10k testing i.e, 12k training 3 val 2.5 test

```
Easy Train counts: {0: 1185, 1: 1349, 2: 1191, 3: 1226, 4: 1169, 5: 108  
4, 6: 1183, 7: 1253, 8: 1170, 9: 1190}
```

```
Easy Val counts : {0: 296, 1: 337, 2: 298, 3: 307, 4: 292, 5: 271, 6: 2  
96, 7: 313, 8: 293, 9: 297}
```

```
Hard Test counts : {0: 245, 1: 284, 2: 258, 3: 253, 4: 245, 5: 223, 6: 2  
40, 7: 257, 8: 243, 9: 252}
```

For fast Training results I used GPU-T4 in kaagle, NUM WORKERS=2
and PIN MEMORY = True in batch loaders in PYTORCH framework



SIMPLE CNN:

```
class LazyCNN(nn.Module):
    def __init__(self):
        super().__init__()

        self.conv1 = nn.Conv2d(3, 6, kernel_size=9, stride=2, padding=4) # 28 → 14
        self.conv2 = nn.Conv2d(6, 8, kernel_size=7, stride=2, padding=3) # 14 → 7
        self.conv3 = nn.Conv2d(8, 16, kernel_size=3, stride=1, padding=1) # 7 → 7

        self.fc = nn.Linear(16 * 7 * 7, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = F.relu(self.conv2(x))
        x = F.relu(self.conv3(x))
        x = torch.flatten(x, 1)
        return self.fc(x)

transform = transforms.ToTensor()
```

Training Accuracy:95.15%

Easy Validation Accuracy: 95.13%

HARD test accuracy :7.04%

RESNET 18:

Ya I didn't give up when my first attempt was hard test accuracy:90 without improper data transformations according to pytorch documentation and I didn't unfreeze the models conv's

```
resnet_transform = transforms.Compose([
    transforms.Resize(256, interpolation=InterpolationMode.BILINEAR),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(
        mean=[0.485, 0.456, 0.406],
        std=[0.229, 0.224, 0.225]
    )
])
```

When unfreezed 2 layers models test accuracy was 0 percentage yes even **HUGE BRAINS ARE biased**

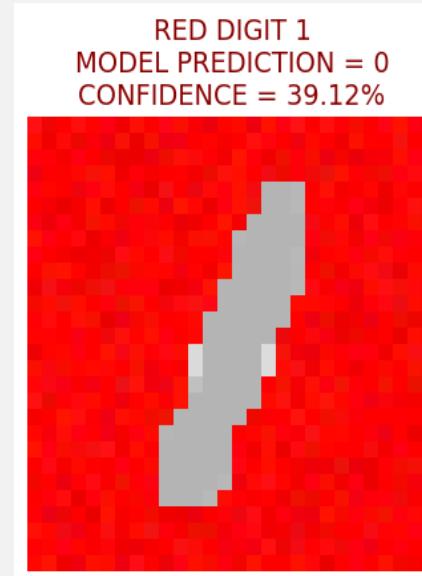
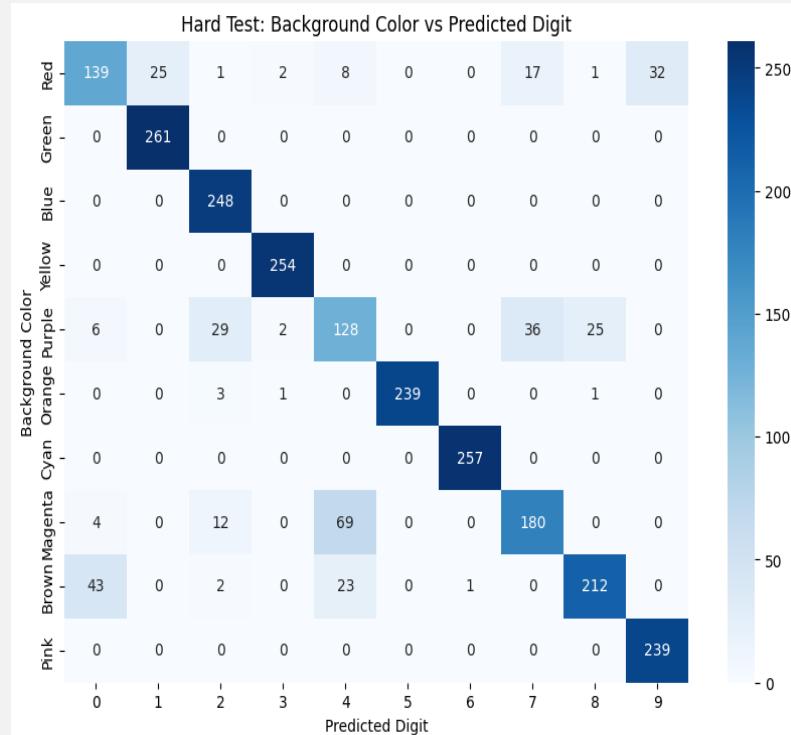
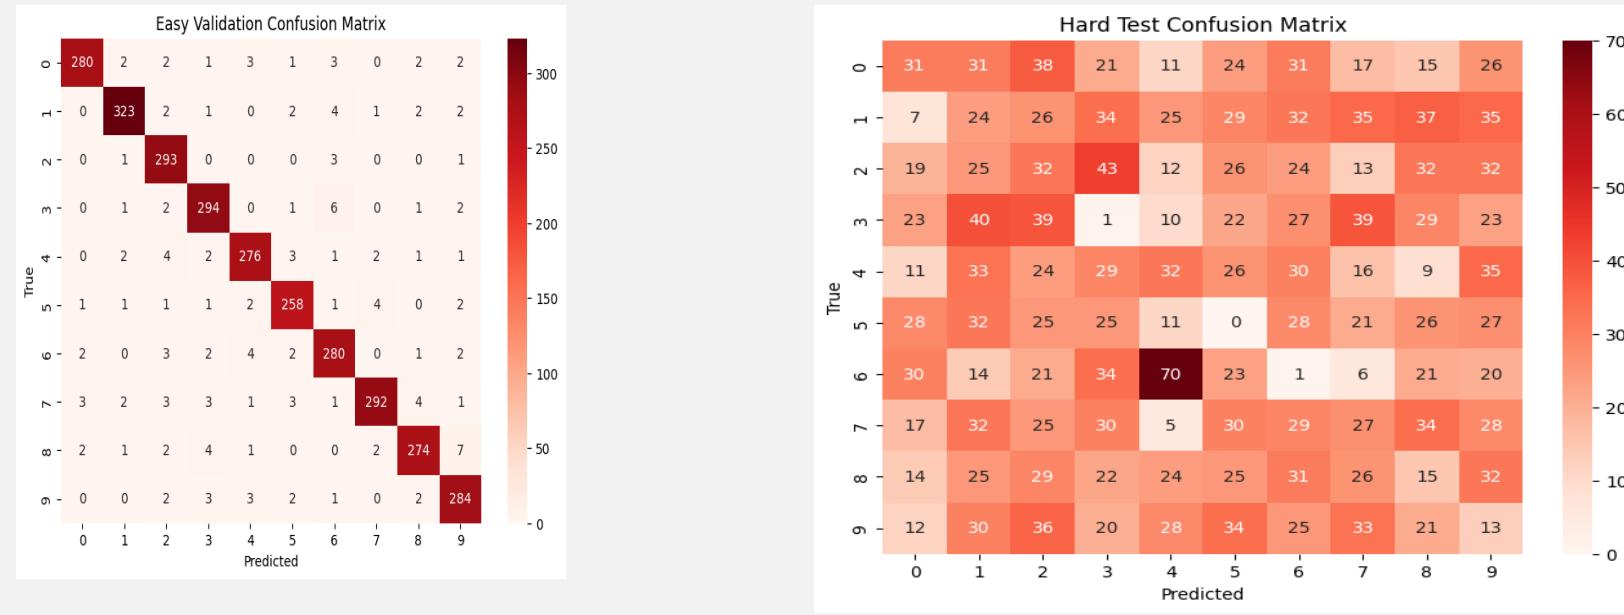
When unfreezed only layer4+fc only:

Training Accuracy:95.36%

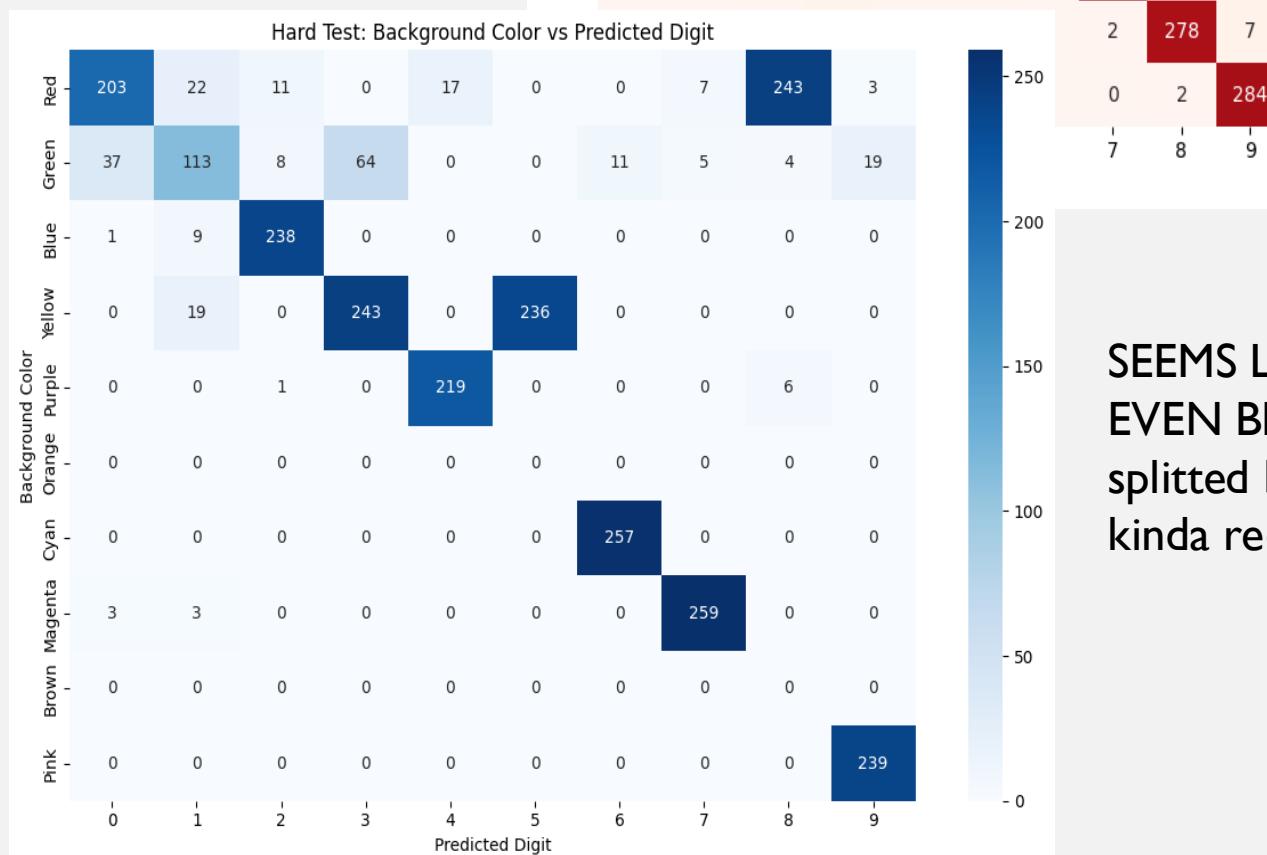
Easy Validation Accuracy: 95.23%

HARD test accuracy :7.56%

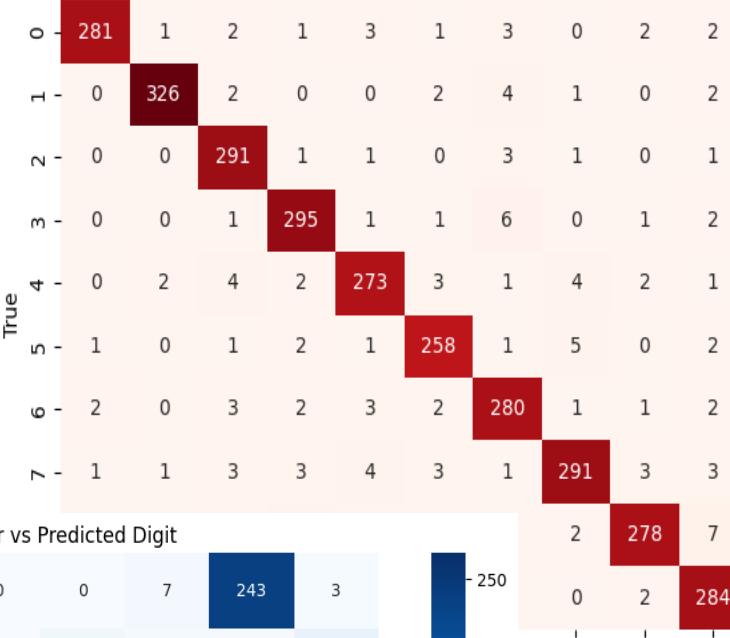
SIMPLE CNN:



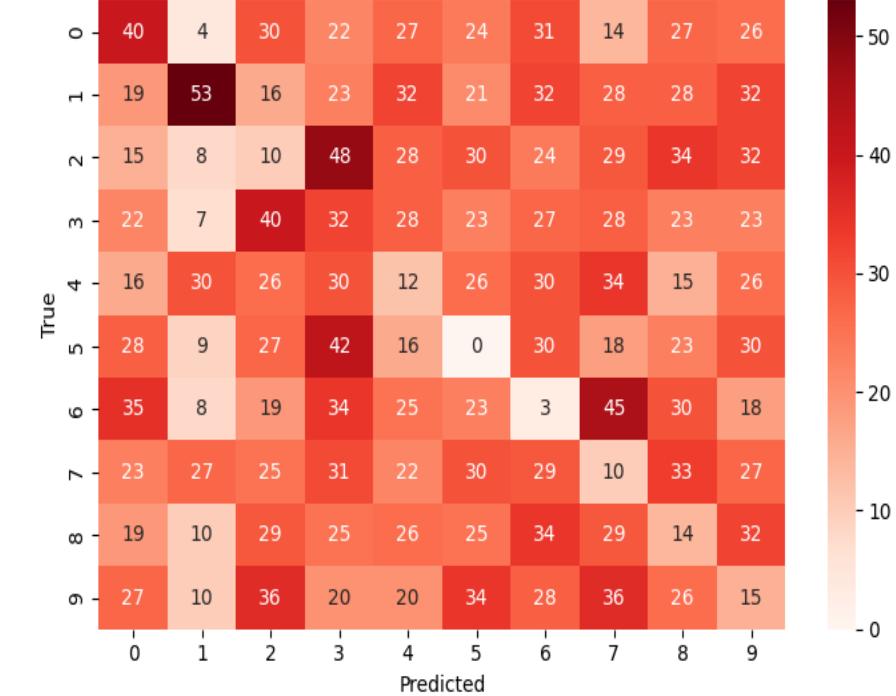
RESNET 18:



Easy Validation Confusion Matrix



Hard Test Confusion Matrix (Fractioned)



SEEMS LIKE RESNET IS EVEN BIASED : red is splitted because brown is kinda reddish brown

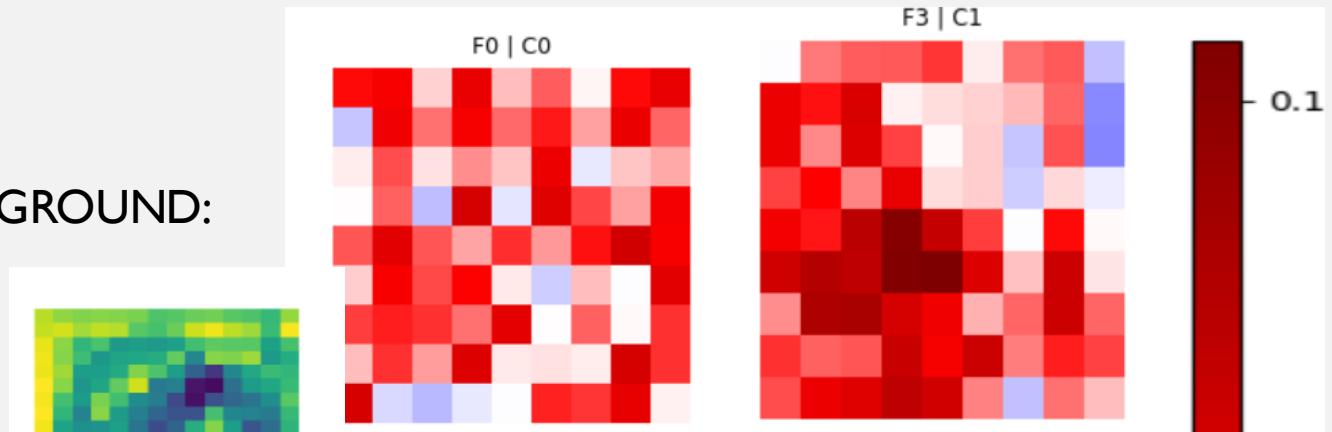


Task 2- The Prober:

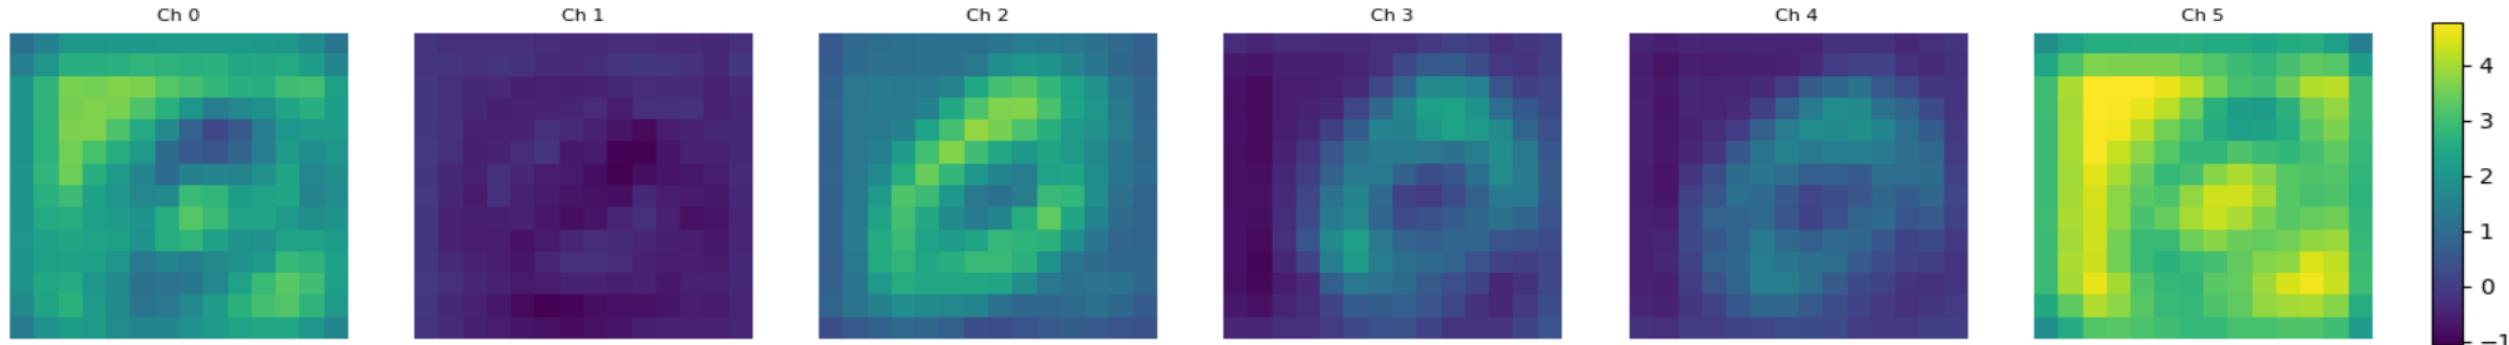
These are simple 3LCNN

FILTERS MOST WEIGHTS ARE FOCUSED ON BACKGROUND:

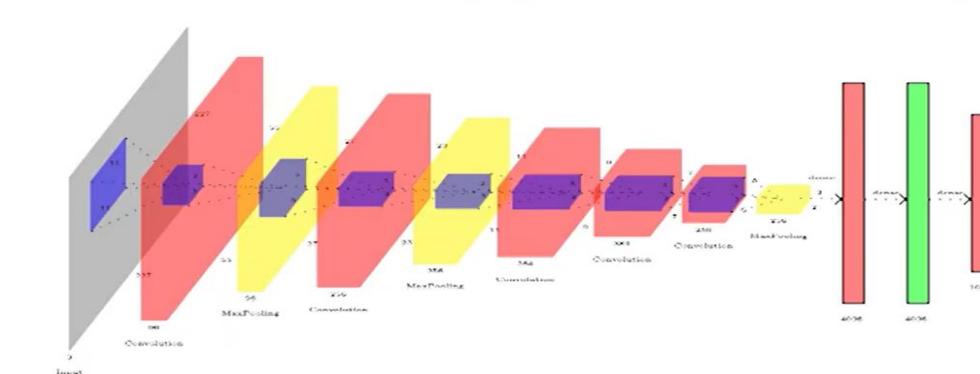
I ensured proper scaling is done for feature maps :



Layer: Conv1 | Range: [-1.07, 4.77]

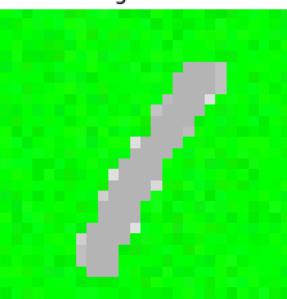


I created image embeddings of 3 Lsimple cnn:

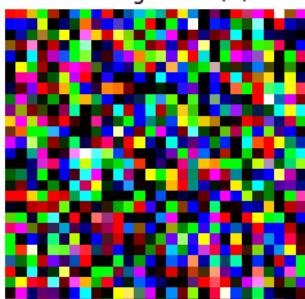


- ϕ_0 : Embedding of an image of interest
 - X : Random image (say zero image)
 - Repeat
 - Forward pass using X and compute $\phi(x)$.
 - Compute
- $$\mathcal{L}(i) = \|\phi(x) - \phi_0\|^2 + \lambda \|\phi(x)\|_6^6$$
- $i_k = i_k - \eta \frac{\partial \mathcal{L}(i)}{\partial i_k}$

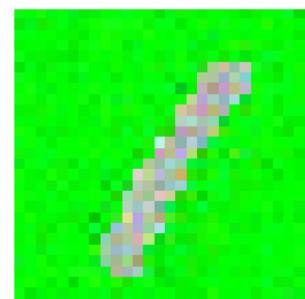
Original '1'



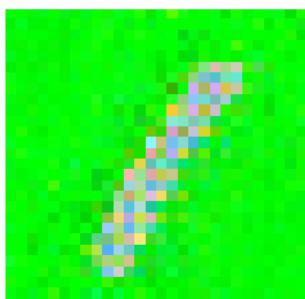
Starting Noise (X)



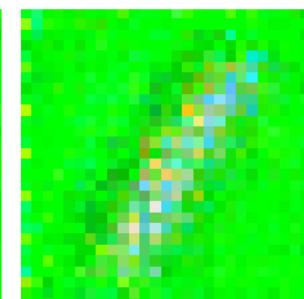
From conv1



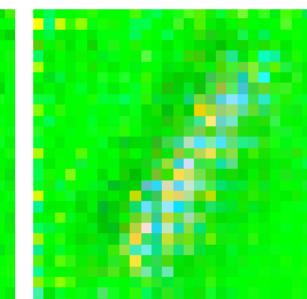
From relu1



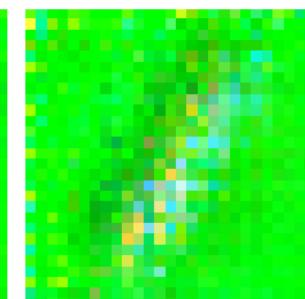
From conv2



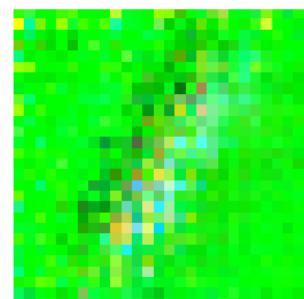
From relu2



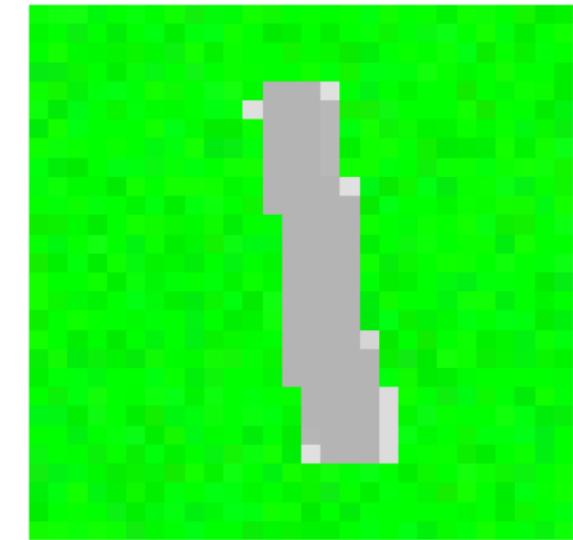
From conv3



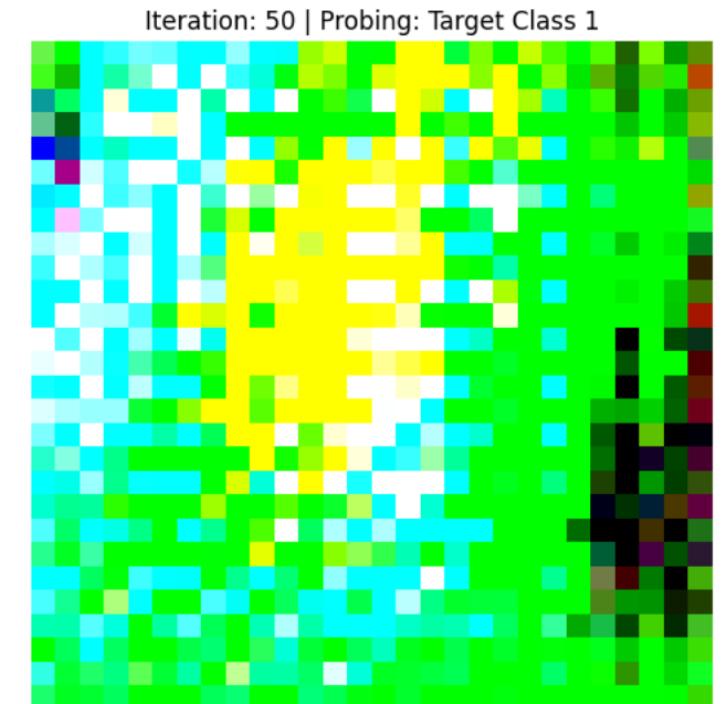
From relu3



I deep dreamed :

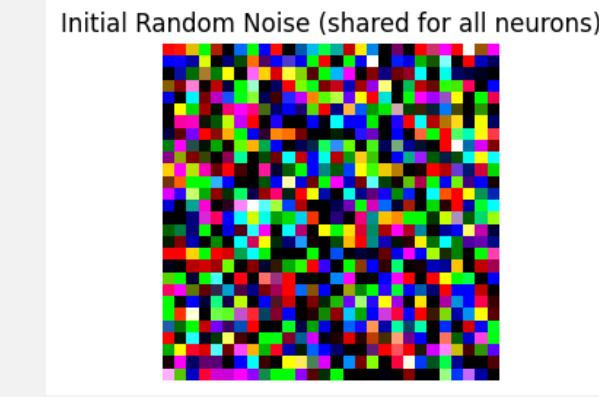


--- Step 2: Precomputing 50 iterations ---
100% | 50/50 [00:00<00:00, 157.17it/s]
Iteration: 50

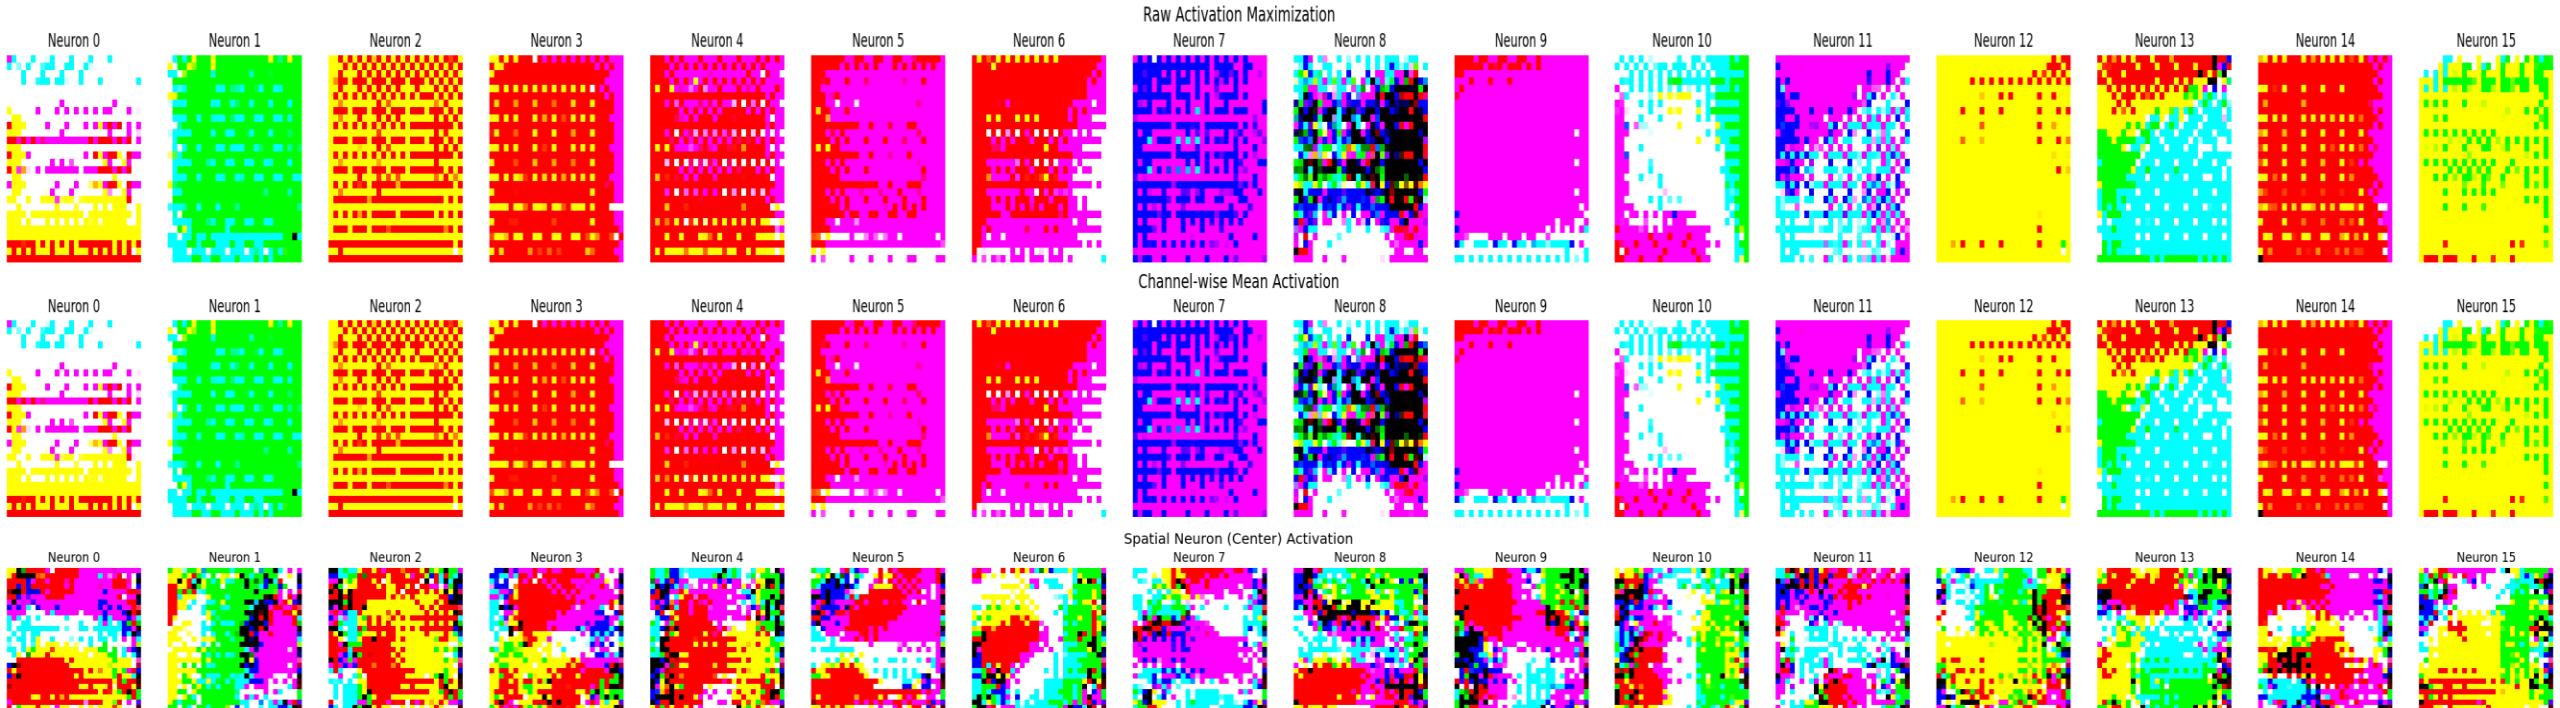


Iteration: 50 | Probing: Target Class 1

Finally! from optimizable image tensor to : raw|channelwise mean|spatial activation
mean activations @ conv3 for simple3LCNN:

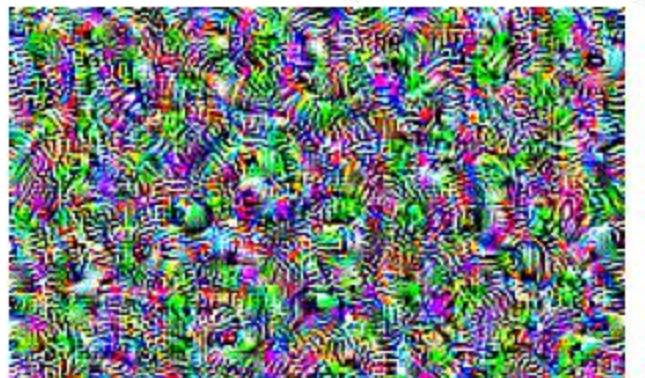


Let's see these ending in task 6! YES THEY ARE COLOURED , TEXTUREDly activated
ARE THEY HAVING DIGIT ACTIVATIONS EVEN?????

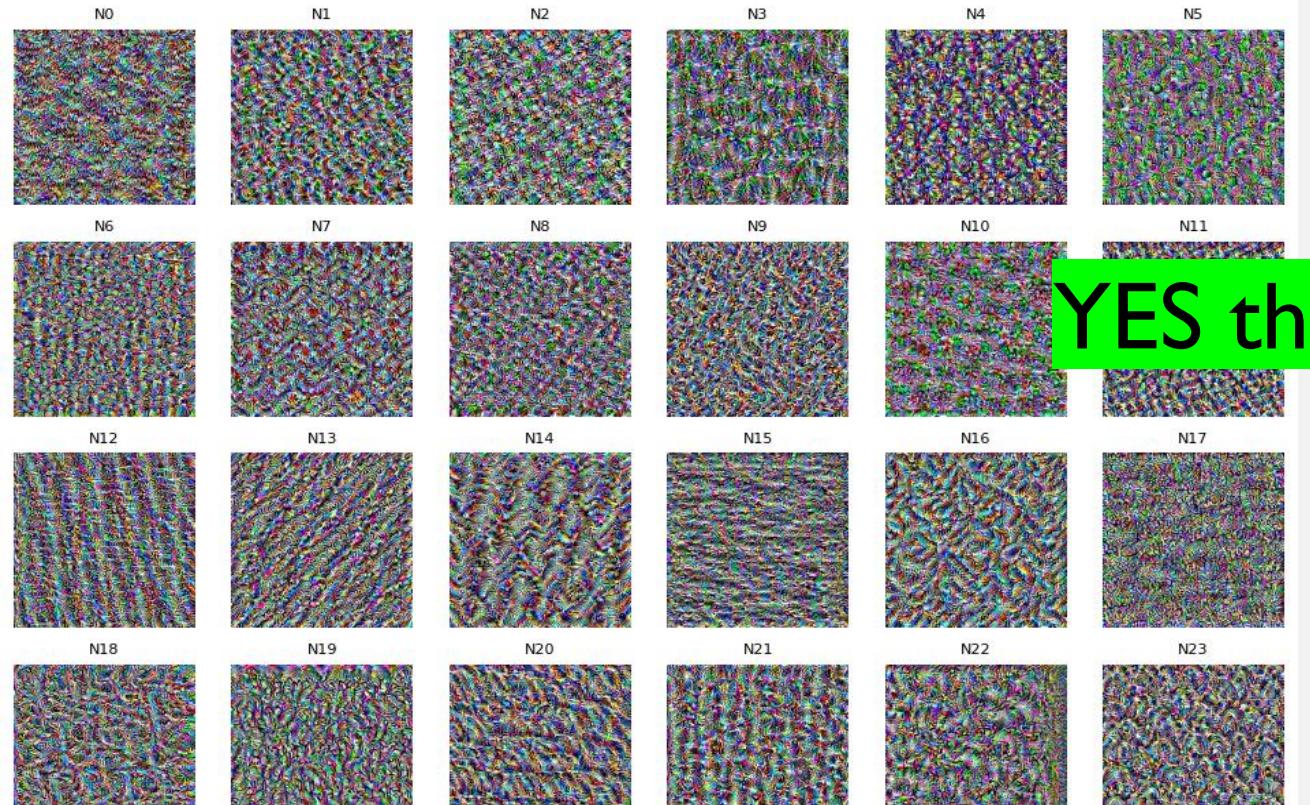


FROM RESNET18:

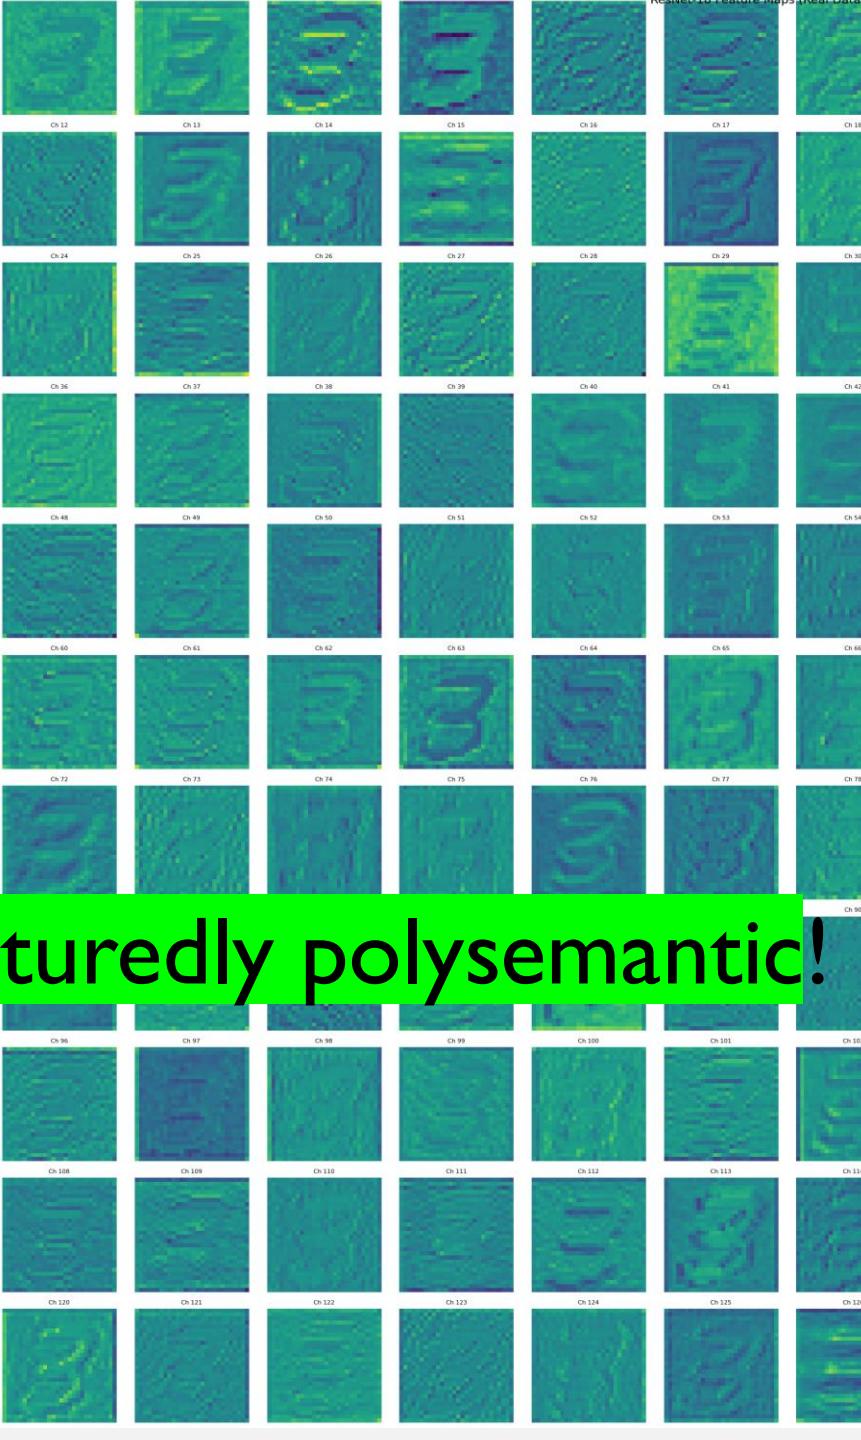
Neuron 5 Preference (Gradient Ascent)



ResNet-18 Polysemantic Neuron Probing (GA from Noise)



Original Dataset Image

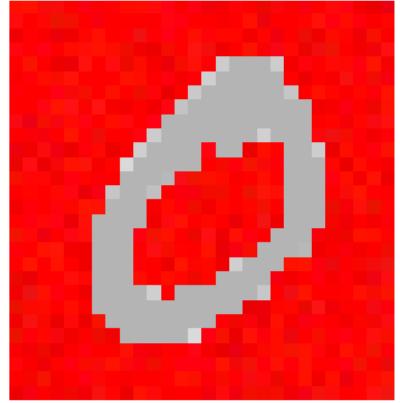


YES they are texturedly polysemantic!

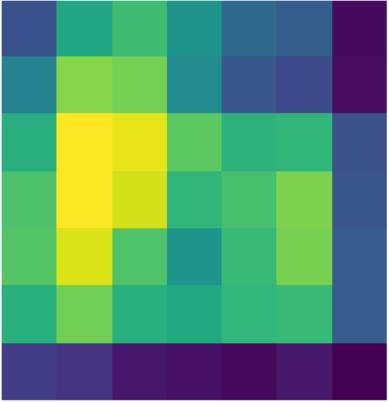
MADE GRAD CAM!!!!-interpolation problem with resnet18:

LazyCNN — Biased example: Red 0 (easy/train)
0_00001.png | pred=0 | target=0 | heatmap MAE=0.033

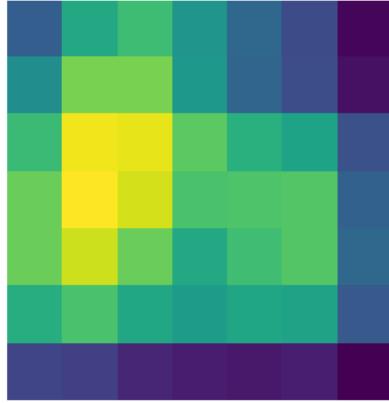
Original



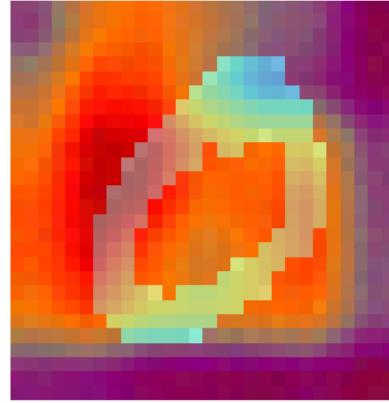
Our heatmap
(viridis, 7x7)



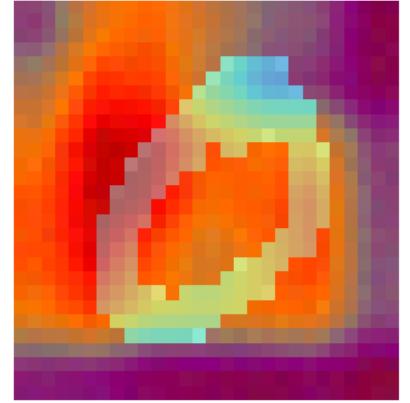
Library heatmap
(viridis, 7x7)



Our overlay
(JET)

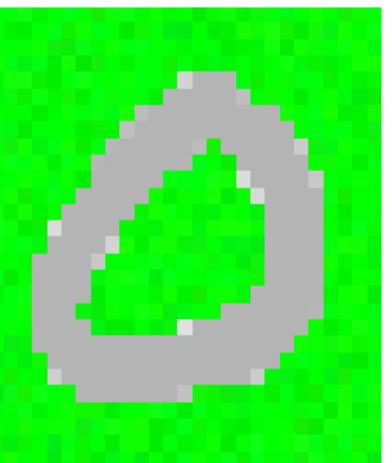


Library overlay
(JET)

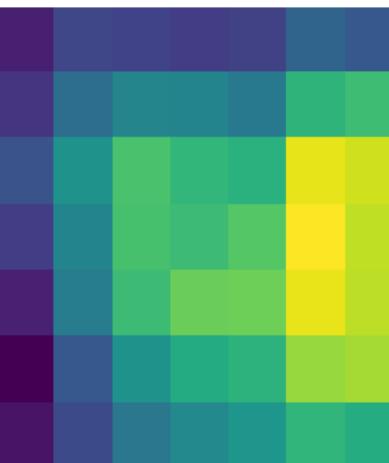


LazyCNN — Conflicting example: Green 0 (hard/test)
0_00296.png | pred=1 | target=1 | heatmap MAE=0.021

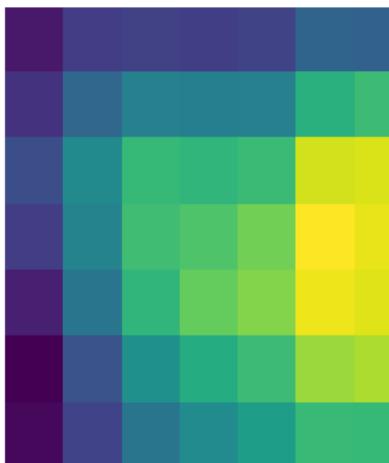
Original



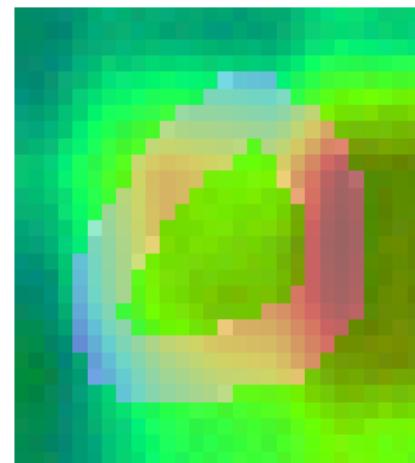
Our heatmap
(viridis, 7x7)



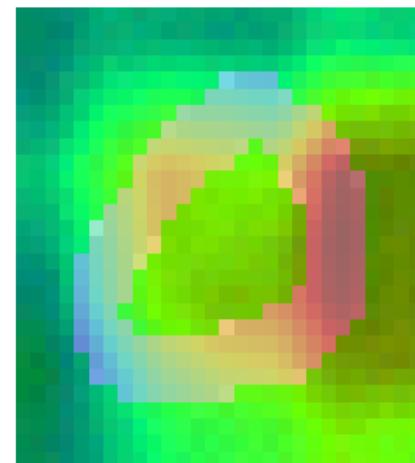
Library heatmap
(viridis, 7x7)



Our overlay
(JET)

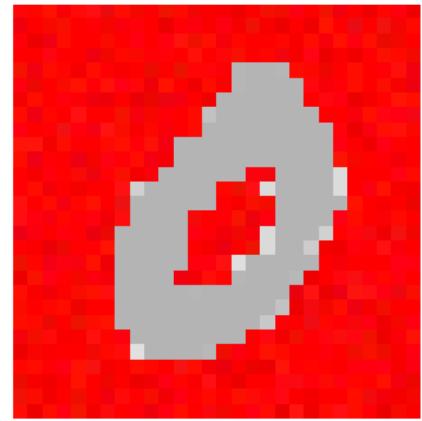


Library overlay
(JET)

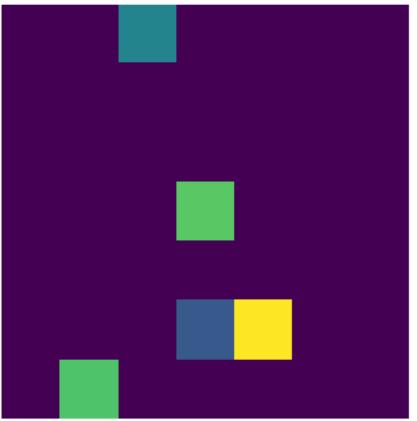


ResNet18 — Biased example: Red 0 (easy/train)
0_00037.png | pred=5 | target=5 | heatmap MAE=0.047

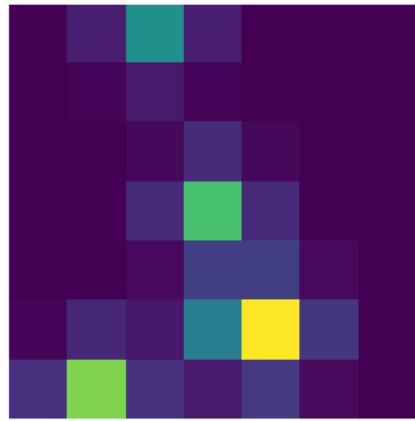
Original



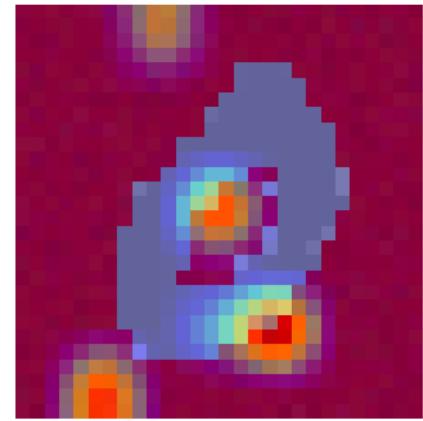
Our heatmap
(viridis, 7x7)



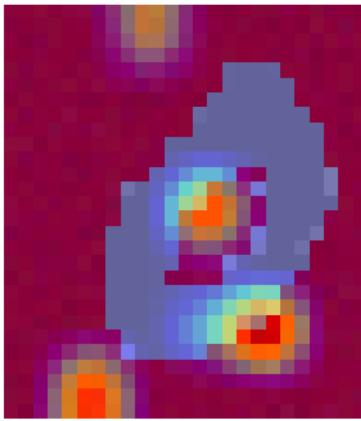
Library heatmap
(viridis, 7x7)



Our overlay
(JET)

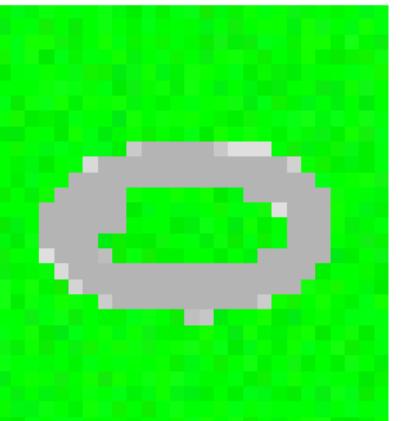


Library overlay
(JET)



ResNet18 — Conflicting example: Green 0 (hard/test)
0_01191.png | pred=3 | target=3 | heatmap MAE=0.029

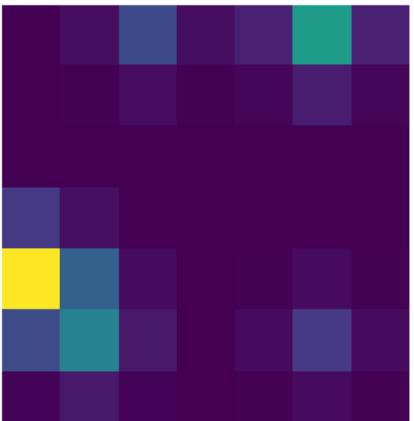
Original



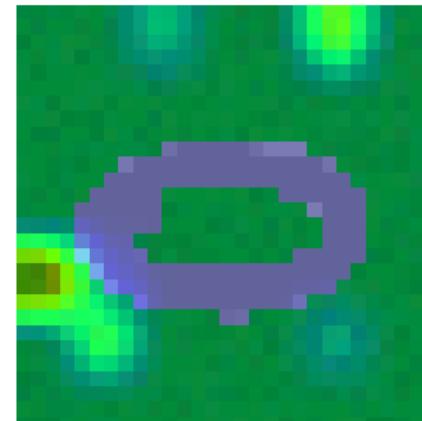
Our heatmap
(viridis, 7x7)



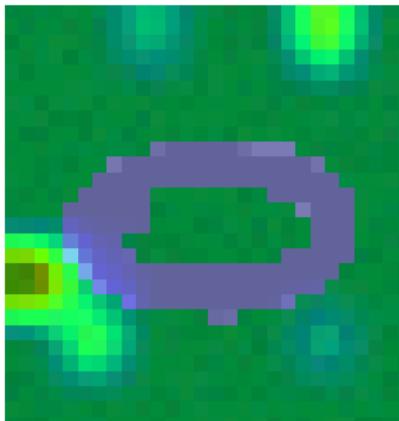
Library heatmap
(viridis, 7x7)



Our overlay
(JET)



Library overlay
(JET)



Task 4 - The Intervention :

My hypo's Method I: USED BATCH NORM for translational invariance
:everything was same transformation.....

| Train Acc: 99.38% | Val Acc: 98.27%

Hard Test Accuracy: 80.84

HAKUNA MATATA!!!!

```
import torch.nn.functional as F

class RobustCNN(nn.Module):
    def __init__(self):
        super().__init__()

        self.conv1 = nn.Sequential(
            nn.Conv2d(3, 32, 5, padding=2),
            nn.BatchNorm2d(32),
            nn.ReLU(),
            nn.MaxPool2d(2) # 28 -> 14
        )

        self.conv2 = nn.Sequential(
            nn.Conv2d(32, 64, 3, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(),
            nn.MaxPool2d(2) # 14 -> 7
        )

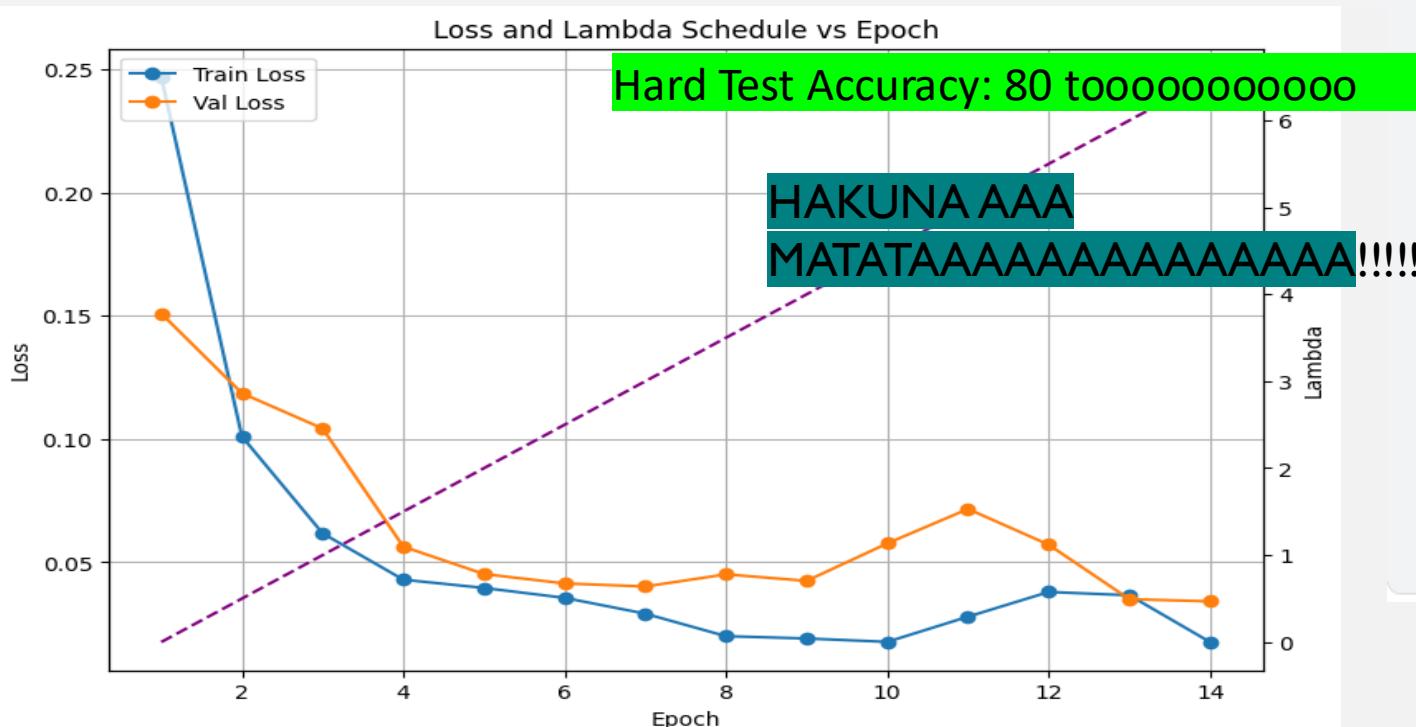
        self.conv3 = nn.Sequential(
            nn.Conv2d(64, 128, 3, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU()
        )

        self.fc = nn.Linear(128 * 7 * 7, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = self.conv2(x)
        x = self.conv3(x)
        x = torch.flatten(x, 1)
        return self.fc(x)
```

Method 2: I USED BATCH NORM + COLOUR PENALISATION with OWN SCHEDULER like lr scheduler for translational invariance and colour invariance
:everything was same transformation

```
def lambda_schedule(epoch):
    return LAMBDA_MAX * (epoch / EPOCHS)
```



```
def color_gradient_penalty(images, logits, labels):
    ce_loss = criterion(logits, labels)

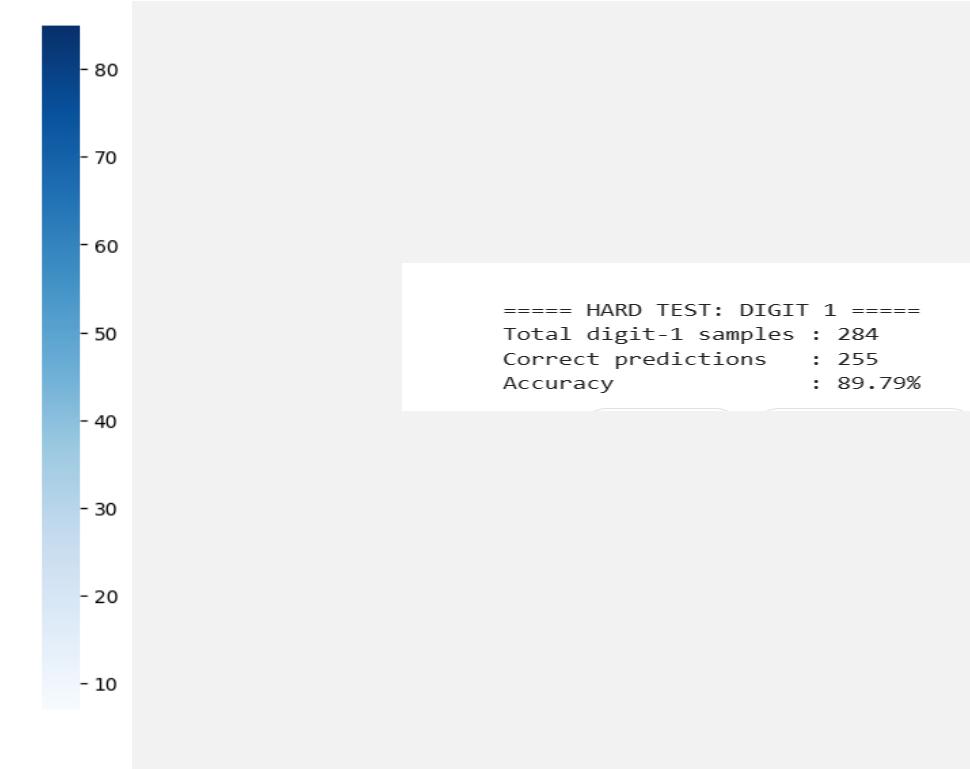
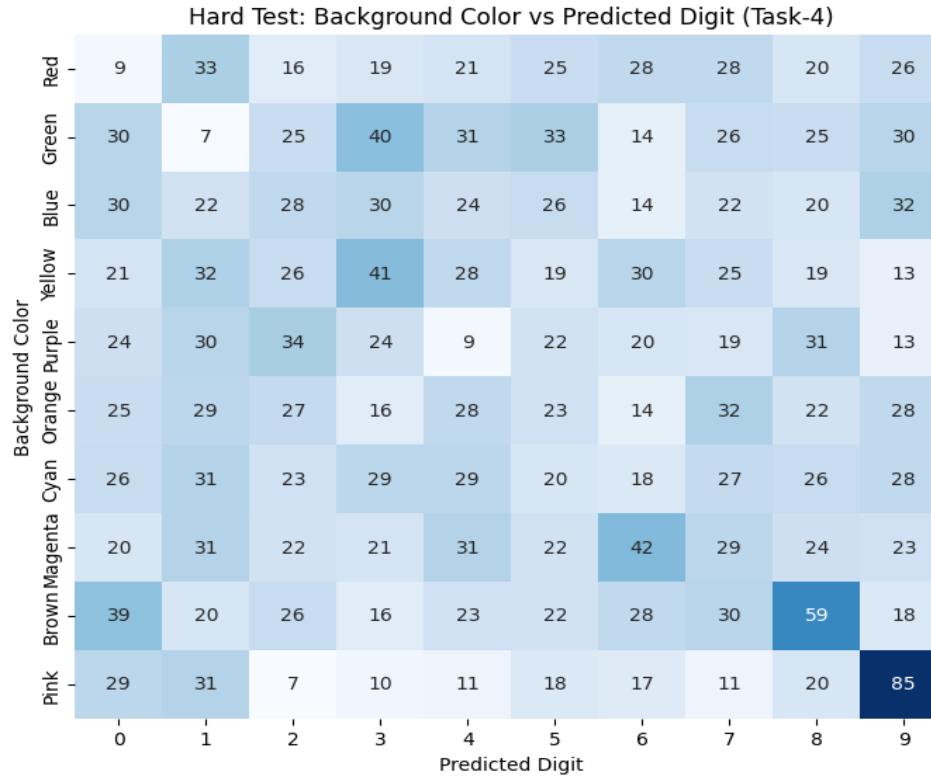
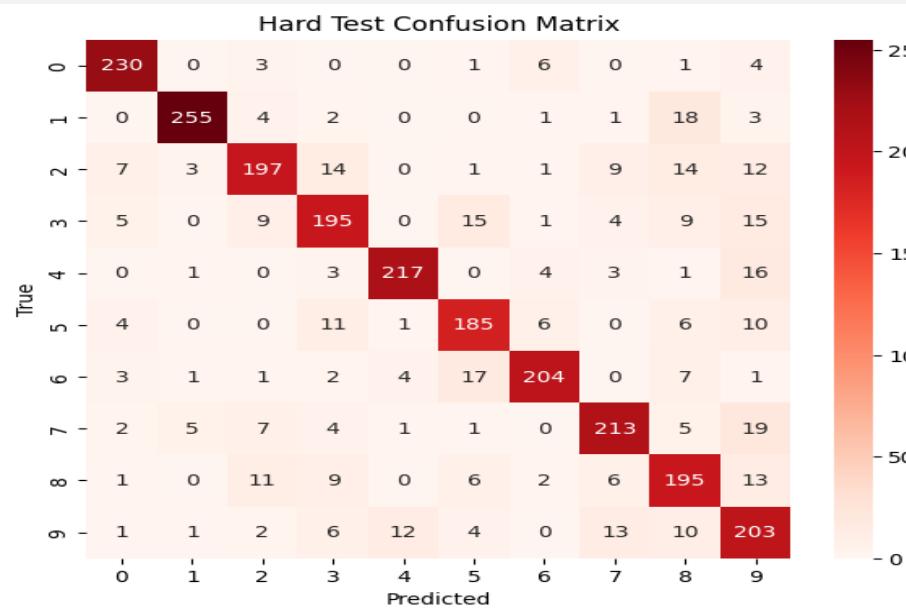
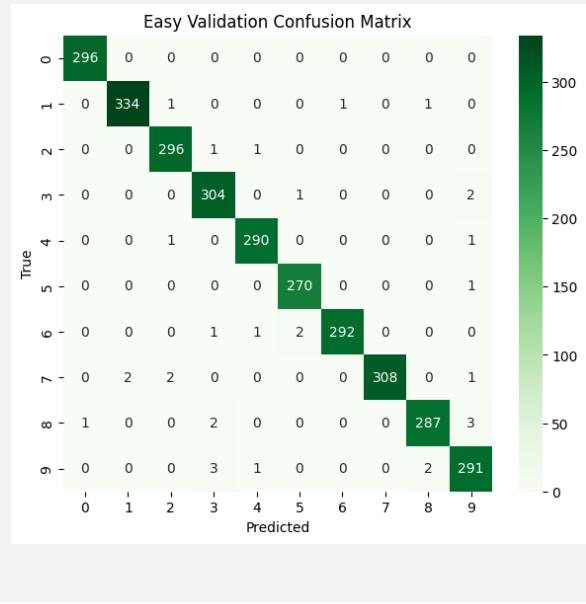
    grads = torch.autograd.grad(
        outputs=logits.gather(1, labels.view(-1, 1)).sum(),
        inputs=images,
        create_graph=True
    )[0]

    grad_r = grads[:, 0]
    grad_g = grads[:, 1]
    grad_b = grads[:, 2]

    color_penalty = torch.var(
        torch.stack([grad_r, grad_g, grad_b], dim=1),
        dim=1
    ).mean()

    return ce_loss, color_penalty
```

TRAIN Acc: 99.79% | Val Acc: 98.93%



Task 5:The Invisible Cloak

: FIRSTLY I was dumb to attack I tried OPTUNA lib –for selection of HYPERPARAMETER's of model uses basyeain optimization in search space,its generalized used in model training ,can be used in Quant task but here I bruteforced to select perturbation of which NOISE : 7 is predicted as 3

No use which I thought for ,I thought it will endup in UNIVERSAL ATTACK like:

The thought is start with noise such that

Noise image := difference of (7 predicted as most confident

3 means there is something making the model as 3
And least confident 7 and 2nd logit as 3) started
That noise and moved to a good noise with optuna for
lazy model so that we might get noise such that
Every image 7 predicted as 3!

Yeah failed because model is lazy such that its
understood what to focus on even –colour pixel by
pixel !

```
x7 = get_clean_7(robust_model, hard_loader)

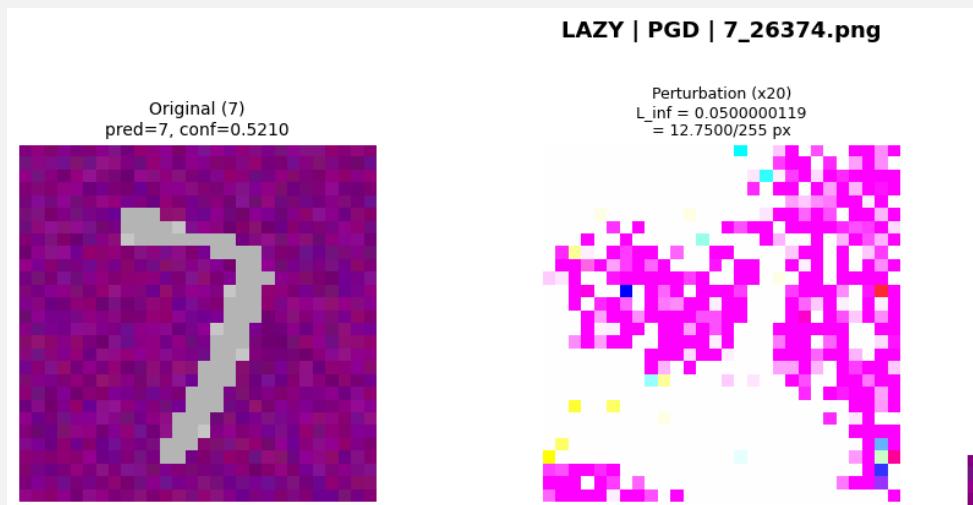
study = optuna.create_study(direction="maximize")
study.optimize(
    lambda trial: optuna_noise_objective(trial, robust_model, x7),
    n_trials=30 # keep small, this is expensive
)

print("Best confidence for class 3:", study.best_value)
```

```
[I 2026-02-03 08:00:40,151] A new study created in memory with name: no-name-c854384c-828c-475c-ae48-8e947dc5fc81
[I 2026-02-03 08:00:40,510] Trial 0 finished with value: 8.954025361163076e-07 and parameters: {'d_0_0_0': -0.006588092390444847, 'd_0_0_1': 0.017820663535946912, 'd_0_0_2': 0.007576114907914287, 'd_0_0_3': -0.03818101061627274, 'd_0_0_4': 0.006504265016708963, 'd_0_0_5': 0.023521812198807196, 'd_0_0_6': -0.0446697436255481, 'd_0_0_7': 0.027499579655146628, 'd_0_0_8': -0.03107002680560903, 'd_0_0_9': -0.017593068687640226, 'd_0_0_10': 0.044264936554160667, 'd_0_0_11': -0.03923716117543733, 'd_0_0_12': -0.02490594024638213, 'd_0_0_13': -0.037772130992795784, 'd_0_0_14': -0.048553117797209014, 'd_0_0_15': 0.04039406694786561, 'd_0_0_16': 0.013801237129546257, 'd_0_0_17': 4.526836480909374e-05, 'd_0_0_18': -0.0375947033426019, 'd_0_0_19': -0.030978499431264054, 'd_0_0_20': -0.0330415433513798, 'd_0_0_21': 0.039906920135984306, 'd_0_0_22': 0.00115547115512856, 'd_0_0_23': 0.021240804260968263, 'd_0_0_24': 0.0030990824435015857, 'd_0_0_25': -0.04085899769015843, 'd_0_0_26': 0.04895600520273734, 'd_0_0_27': 0.03650034669749162, 'd_0_0_1_0': -0.021753804214810404, 'd_0_0_1_1': 0.0007057748131277936, 'd_0_0_1_2': 0.002315329451426097, 'd_0_0_1_3': -0.011055639345919535, 'd_0_0_1_4': -0.010517936146262008, 'd_0_0_1_5': 0.017615763952363778, 'd_0_0_1_6': -0.034324469728518756, 'd_0_0_1_7': -0.025557474084808674, 'd_0_0_1_8': -0.001984357311278148, 'd_0_0_1_9': -0.02138325214696936, 'd_0_0_1_10': -0.04361078169948654, 'd_0_0_1_11': -0.01449950641615633, 'd_0_0_1_12': -0.020458528482044627, 'd_0_0_1_13': 0.021357883141971234, 'd_0_0_1_14': 0.01318325214696936, 'd_0_0_1_15': -0.04361078169948654, 'd_0_0_1_16': 0.01685095195384615, 'd_0_0_1_17': -0.010930881450547744, 'd_0_0_1_18': 0.03401051416516848, 'd_0_0_1_19': 0.004614809528294239, 'd_0_0_1_20': -0.048851253271636225, 'd_0_0_1_21': 0.0062046099380096, 'd_0_0_1_22': 0.04848142969590865, 'd_0_0_1_23': -0.008749893700263134, 'd_0_0_1_24': -0.03794721714309064, 'd_0_0_1_25': -0.03495434668603256, 'd_0_0_1_26': -0.02403623694417584, 'd_0_0_1_27': -0.0015812201599157832, 'd_0_0_2_0': -0.005618714491758811, 'd_0_0_2_1': 0.015676336035307242, 'd_0_0_2_2': -0.02319188579670174, 'd_0_0_2_3': 0.013486602274012188, 'd_0_0_2_4': -0.01487315272496701, 'd_0_0_2_5': -0.049323785012075176, 'd_0_0_2_6': 0.003399627762036915, 'd_0_0_2_7': -0.02256699940179101, 'd_0_0_2_8': 0.016475657339745403, 'd_0_0_2_9': -0.028027877831922222, 'd_0_0_2_10': 0.024350010370076322, 'd_0_0_2_11': -0.016254747621216162, 'd_0_0_2_12': 0.04875845960410037, 'd_0_0_2_13': 0.04887796032795677, 'd_0_0_2_14': -0.03718591223993086, 'd_0_0_2_15': -0.02109145865491624, 'd_0_0_2_16': 0.04949908852780874, 'd_0_0_2_17': 0.002709690153261314, 'd_0_0_2_18': 0.04106131791512811, 'd_0_0_2_19': 0.0118783117378235, 'd_0_0_2_20': -0.004879442297517987, 'd_0_0_2_21': 0.010973296014053668, 'd_0_0_2_22': -0.0489869000946824, 'd_0_0_2_23': 0.031951996666598315, 'd_0_0_2_24': 0.04468638415190866, 'd_0_0_2_25': -0.03919985217357931, 'd_0_0_2_26': -0.03912652389236371, 'd_0_0_2_27': 0.0452900500234126, 'd_0_0_2_28': -0.024174955340581495, 'd_0_0_2_29': 0.0401566148025776, 'd_0_0_2_30': 0.048119050744989844, 'd_0_0_2_31': 0.02088807175594694, 'd_0_0_2_32': 0.006058920065387177, 'd_0_0_2_33': 0.0452900500234126, 'd_0_0_2_34': 0.0244053505418814, 'd_0_0_2_35': 0.011719966449881526, 'd_0_0_2_36': 0.028218177009384407, 'd_0_0_2_37': 0.007443262416216581, 'd_0_0_2_38': -0.0030752338215108907, 'd_0_0_2_39': 0.03530637664965951, 'd_0_0_2_40': -0.0022822042054352135, 'd_0_0_2_41': -0.02815510155195199, 'd_0_0_2_42': -0.0015184139651492554, 'd_0_0_2_43': -0.013314959853501021, 'd_0_0_2_44': 0.003504432230112667, 'd_0_0_2_45': 0.02991494971499055, 'd_0_0_2_46': 0.0023834148898164015, 'd_0_0_2_47': -0.012183921912513343, 'd_0_0_2_48': -0.03976204351868539, 'd_0_0_2_49': -0.047067037707233, 'd_0_0_2_50': 0.025529937170663805, 'd_0_0_2_51': -0.0016573785481718492, 'd_0_0_2_52': 0.04486027565262524, 'd_0_0_2_53': -0.006495494952256063, 'd_0_0_2_54': -0.03791299839588485, 'd_0_0_2_55': 0.0072465369984772016, 'd_0_0_2_56': 0.04802508231114855, 'd_0_0_2_57': -0.02078182162524841, 'd_0_0_2_58': -0.012638428278296913, 'd_0_0_2_59': -0.042631494253819225, 'd_0_0_2_60': 0.011781148747205782, 'd_0_0_2_61': -0.03205506740774111, 'd_0_0_2_62': 0.024054392349926262, 'd_0_0_2_63': -0.040441818360260035, 'd_0_0_2_64': -0.0034200199318815266, 'd_0_0_2_65': -0.01510627039043129, 'd_0_0_2_66': 0.0072896458661871125, 'd_0_0_2_67': -0.024592887423311118, 'd_0_0_2_68': -0.004273117795534721, 'd_0_0_2_69': -0.004257113611486642, 'd_0_0_2_70': -0.034156173338392076, 'd_0_0_2_71': 0.04143826576345924, 'd_0_0_2_72': 0.034230861233901, 'd_0_0_2_73': 0.045667853606460965, 'd_0_0_2_74': 0.0490789442504185, 'd_0_0_2_75': 0.04585072429324559, 'd_0_0_2_76': 0.026712438928862672, 'd_0_0_2_77': -0.02078511798940516, 'd_0_0_2_78': -0.010084638045584036, 'd_0_0_2_79': 0.0272304733097752, 'd_0_0_2_80': -0.027250501992987966, 'd_0_0_2_81': 0.0212843519504681, 'd_0_0_2_82': -0.036759686707342566, 'd_0_0_2_83': 0.0377312322986925, 'd_0_0_2_84': 0.0377312322986925, 'd_0_0_2_85': 0.02785037597441274, 'd_0_0_2_86': -0.03920197974362875, 'd_0_0_2_87': -0.048027553230873175, 'd_0_0_2_88': -0.0367111883339759, 'd_0_0_2_89': -0.014269287720923805, 'd_0_0_2_90': -0.027320330918621096, 'd_0_0_2_91': -0.034608466143597284, 'd_0_0_2_92': 0.016236863866469914, 'd_0_0_2_93': 0.03630927565897889, 'd_0_0_2_94': -0.03389313285533424, 'd_0_0_2_95': 0.0057303995278603, 'd_0_0_2_96': -0.004491236684922774, 'd_0_0_2_97': 0.003402259886465163, 'd_0_0_2_98': 0.03356173625890653, 'd_0_0_2_99': 0.0137369604518292, 'd_0_0_3_00': 0.021348561703713534, 'd_0_0_3_01': 0.011527964523556446, 'd_0_0_3_02': -0.018132675971463057, 'd_0_0_3_03': 0.02755237879158608, 'd_0_0_3_04': 0.008055281988621441, 'd_0_0_3_05': -0.04738757922197593, 'd_0_0_3_06': 0.04937511021347674, 'd_0_0_3_07': -0.0343586888725052, 'd_0_0_3_08': 0.007417051883644047, 'd_0_0_3_09': 0.0309025187487732, 'd_0_0_3_10': -0.0454112109188, 'd_0_0_3_11': 0.03399781329843274, 'd_0_0_3_12': 0.049447007570451596, 'd_0_0_3_13': 0.0201770082020890176, 'd_0_0_3_14': 0.03836673051460641, 'd_0_0_3_15': 0.009114908899271182, 'd_0_0_3_16': 0.008244722823564796, 'd_0_0_3_17': 0.015467863231927684, 'd_0_0_3_18': 0.03567767742116139, 'd_0_0_3_19': 0.023092353202905594, 'd_0_0_3_20': -0.036840022081690176, 'd_0_0_3_21': 0.009114908899271182, 'd_0_0_3_22': -0.036840022081690176, 'd_0_0_3_23': 0.03836673051460641, 'd_0_0_3_24': 0.009114908899271182, 'd_0_0_3_25': -0.036840022081690176, 'd_0_0_3_26': 0.03567767742116139, 'd_0_0_3_27': -0.036840022081690176, 'd_0_0_3_28': 0.009114908899271182, 'd_0_0_3_29': -0.036840022081690176, 'd_0_0_3_30': 0.03567767742116139, 'd_0_0_3_31': -0.036840022081690176, 'd_0_0_3_32': 0.009114908899271182, 'd_0_0_3_33': -0.036840022081690176, 'd_0_0_3_34': 0.03567767742116139, 'd_0_0_3_35': -0.036840022081690176, 'd_0_0_3_36': 0.009114908899271182, 'd_0_0_3_37': -0.036840022081690176, 'd_0_0_3_38': 0.03567767742116139, 'd_0_0_3_39': -0.036840022081690176, 'd_0_0_3_40': 0.009114908899271182, 'd_0_0_3_41': -0.036840022081690176, 'd_0_0_3_42': 0.03567767742116139, 'd_0_0_3_43': -0.036840022081690176, 'd_0_0_3_44': 0.009114908899271182, 'd_0_0_3_45': -0.036840022081690176, 'd_0_0_3_46': 0.03567767742116139, 'd_0_0_3_47': -0.036840022081690176, 'd_0_0_3_48': 0.009114908899271182, 'd_0_0_3_49': -0.036840022081690176, 'd_0_0_3_50': 0.03567767742116139, 'd_0_0_3_51': -0.036840022081690176, 'd_0_0_3_52': 0.009114908899271182, 'd_0_0_3_53': -0.036840022081690176, 'd_0_0_3_54': 0.03567767742116139, 'd_0_0_3_55': -0.036840022081690176, 'd_0_0_3_56': 0.009114908899271182, 'd_0_0_3_57': -0.036840022081690176, 'd_0_0_3_58': 0.03567767742116139, 'd_0_0_3_59': -0.036840022081690176, 'd_0_0_3_60': 0.009114908899271182, 'd_0_0_3_61': -0.036840022081690176, 'd_0_0_3_62': 0.03567767742116139, 'd_0_0_3_63': -0.036840022081690176, 'd_0_0_3_64': 0.009114908899271182, 'd_0_0_3_65': -0.036840022081690176, 'd_0_0_3_66': 0.03567767742116139, 'd_0_0_3_67': -0.036840022081690176, 'd_0_0_3_68': 0.009114908899271182, 'd_0_0_3_69': -0.036840022081690176, 'd_0_0_3_70': 0.03567767742116139, 'd_0_0_3_71': -0.036840022081690176, 'd_0_0_3_72': 0.009114908899271182, 'd_0_0_3_73': -0.036840022081690176, 'd_0_0_3_74': 0.03567767742116139, 'd_0_0_3_75': -0.036840022081690176, 'd_0_0_3_76': 0.009114908899271182, 'd_0_0_3_77': -0.036840022081690176, 'd_0_0_3_78': 0.03567767742116139, 'd_0_0_3_79': -0.036840022081690176, 'd_0_0_3_80': 0.009114908899271182, 'd_0_0_3_81': -0.036840022081690176, 'd_0_0_3_82': 0.03567767742116139, 'd_0_0_3_83': -0.036840022081690176, 'd_0_0_3_84': 0.009114908899271182, 'd_0_0_3_85': -0.036840022081690176, 'd_0_0_3_86': 0.03567767742116139, 'd_0_0_3_87': -0.036840022081690176, 'd_0_0_3_88': 0.009114908899271182, 'd_0_0_3_89': -0.036840022081690176, 'd_0_0_3_90': 0.03567767742116139, 'd_0_0_3_91': -0.036840022081690176, 'd_0_0_3_92': 0.009114908899271182, 'd_0_0_3_93': -0.036840022081690176, 'd_0_0_3_94': 0.03567767742116139, 'd_0_0_3_95': -0.036840022081690176, 'd_0_0_3_96': 0.009114908899271182, 'd_0_0_3_97': -0.036840022081690176, 'd_0_0_3_98': 0.03567767742116139, 'd_0_0_3_99': -0.036840022081690176, 'd_0_0_4_00': 0.009114908899271182, 'd_0_0_4_01': -0.036840022081690176, 'd_0_0_4_02': 0.03567767742116139, 'd_0_0_4_03': -0.036840022081690176, 'd_0_0_4_04': 0.009114908899271182, 'd_0_0_4_05': -0.036840022081690176, 'd_0_0_4_06': 0.03567767742116139, 'd_0_0_4_07': -0.036840022081690176, 'd_0_0_4_08': 0.009114908899271182, 'd_0_0_4_09': -0.036840022081690176, 'd_0_0_4_10': 0.03567767742116139, 'd_0_0_4_11': -0.036840022081690176, 'd_0_0_4_12': 0.009114908899271182, 'd_0_0_4_13': -0.036840022081690176, 'd_0_0_4_14': 0.009114908899271182, 'd_0_0_4_15': -0.036840022081690176, 'd_0_0_4_16': 0.009114908899271182, 'd_0_0_4_17': -0.036840022081690176, 'd_0_0_4_18': 0.009114908899271182, 'd_0_0_4_19': -0.036840022081690176, 'd_0_0_4_20': 0.009114908899271182, 'd_0_0_4_21': -0.036840022081690176, 'd_0_0_4_22': 0.009114908899271182, 'd_0_0_4_23': -0.036840022081690176, 'd_0_0_4_24': 0.009114908899271182, 'd_0_0_4_25': -0.036840022081690176, 'd_0_0_4_26': 0.009114908899271182, 'd_0_0_4_27': -0.036840022081690176, 'd_0_0_4_28': 0.009114908899271182, 'd_0_0_4_29': -0.036840022081690176, 'd_0_0_4_30': 0.009114908899271182, 'd_0_0_4_31': -0.036840022081690176, 'd_0_0_4_32': 0.009114908899271182, 'd_0_0_4_33': -0.036840022081690176, 'd_0_0_4_34': 0.009114908899271182, 'd_0_0_4_35': -0.036840022081690176, 'd_0_0_4_36': 0.009114908899271182, 'd_0_0_4_37': -0.036840022081690176, 'd_0_0_4_38': 0.009114908899271182, 'd_0_0_4_39': -0.036840022081690176, 'd_0_0_4_40': 0.009114908899271182, 'd_0_0_4_41': -0.036840022081690176, 'd_0_0_4_42': 0.009114908899271182, 'd_0_0_4_43': -0.036840022081690176, 'd_0_0_4_44': 0.009114908899271182, 'd_0_0_4_45': -0.036840022081690176, 'd_0_0_4_46': 0.009114908899271182, 'd_0_0_4_47': -0.036840022081690176, 'd_0_0_4_48': 0.009114908899271182, 'd_0_0_4_49': -0.036840022081690176, 'd_0_0_4_50': 0.009114908899271182, 'd_0_0_4_51': -0.036840022081690176, 'd_0_0_4_52': 0.009114908899271182, 'd_0_0_4_53': -0.036840022081690176, 'd_0_0_4_54': 0.009114908899271182, 'd_0_0_4_55': -0.036840022081690176, 'd_0_0_4_56': 0.009114908899271182
```

FINALLY ATTACKED! TARGETEDLLL!!!!: with pgd ,tgds,deepfool

up on selecting top5 images of both models predict 7 with 2nd-choice=3: before THAT I GAVE UP AND THOUGHT FINALLY ,ITS NOT OVER UNTIL I WIN !



Top-5 selected images:						
#	p(3)_lazy	p(3)_robust	conf_lazy	conf_robust	file	
1	0.019028	0.010362	0.9492	0.9845	7_31850.png	
		0.009425	0.5210	0.9875	7_26374.png	
		0.006559	0.9557	0.9913	7_23946.png	
		0.010650	0.9954	0.9880	7_06050.png	
		0.037047	0.9962	0.9511	7_18322.png	

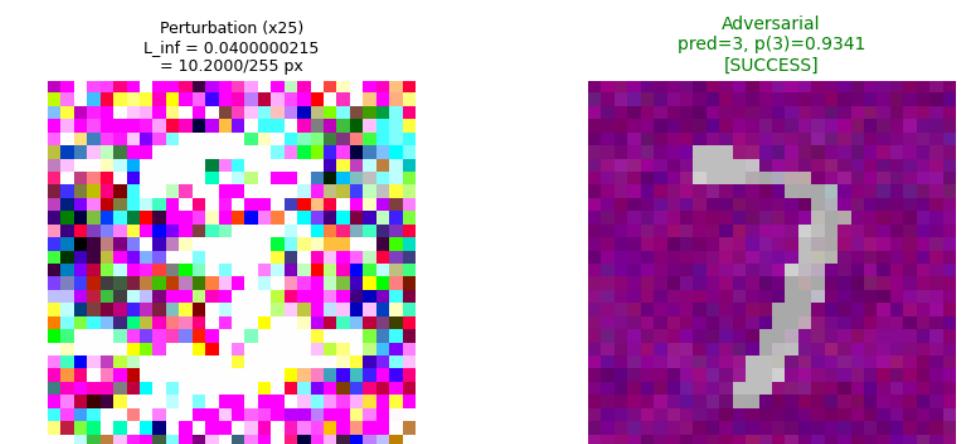
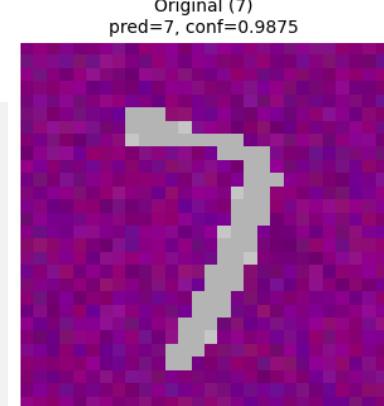
Total 7-images: 6265
Both predict 7 with 2nd-choice=3: 352

Adversarial pred=3, p(3)=0.4985 [FAILED]

YESSSSS
LETS GO::: $L_{\inf} = 0.040000021$

ROBUST | PGD | 7_26374.png

A 28x28 pixel image showing a handwritten digit '7' on a dark purple background. A large, highly colored cluster of pixels (pink, yellow, green, blue) is centered over the digit, representing a more significant perturbation than the original.



=====

SUMMARY: Targeted Attack 7 -> 3 | eps < 0.05 | p(3) > 90%

=====

#	Image	LAZY best	LAZY p(3)	LAZY ok	ROBUST best	ROB p(3)	ROB ok
1	7_31850.png	PGD	0.1998	FAIL	PGD	0.9461	PASS
2	7_26374.png	PGD	0.4985	FAIL	PGD	0.9341	PASS
3	7_23946.png	PGD	0.2904	FAIL	PGD	0.9252	PASS
4	7_06050.png	PGD	0.0817	FAIL	PGD	0.9465	PASS
5	7_18322.png	PGD	0.0620	FAIL	PGD	0.9457	PASS
TOTAL SUCCESSFUL		0/5		5/5			

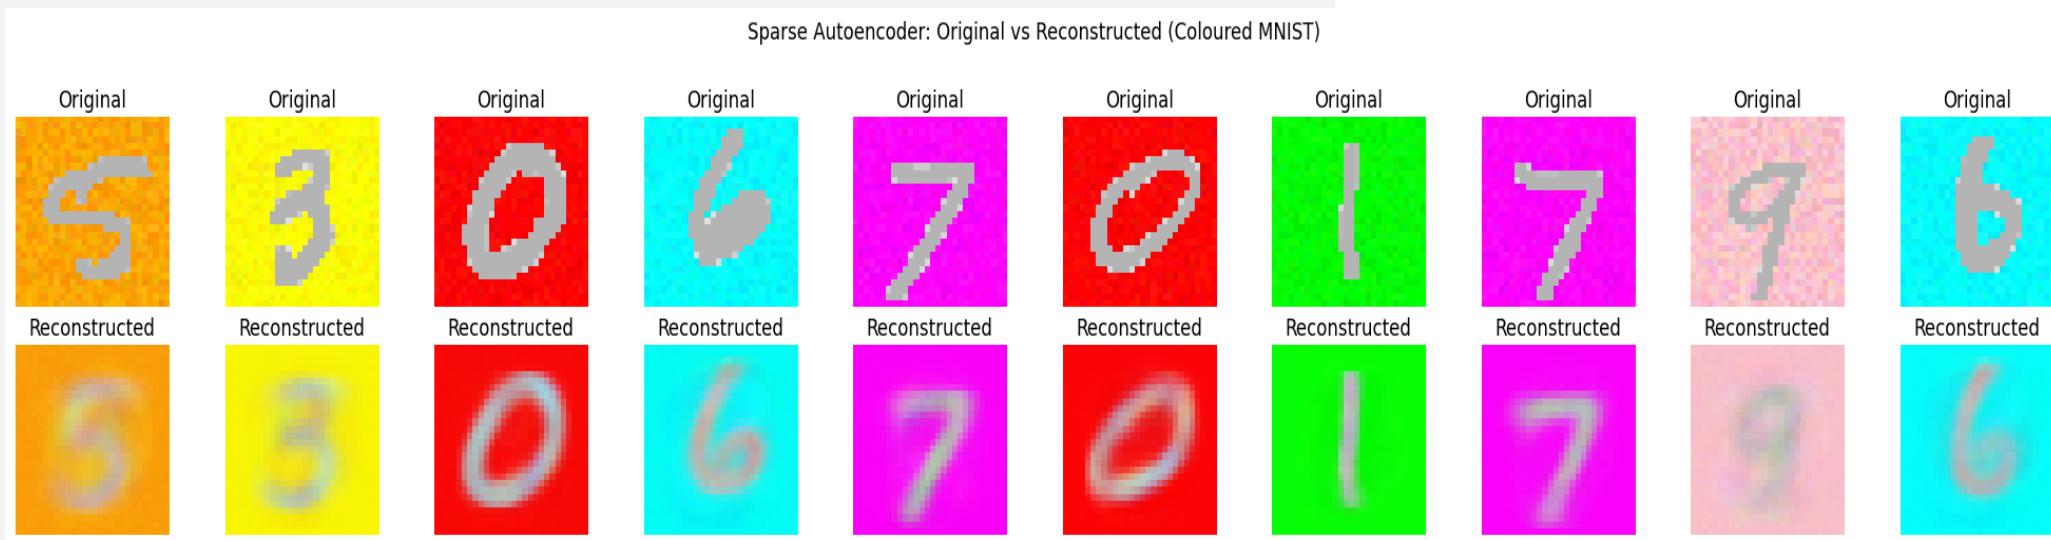
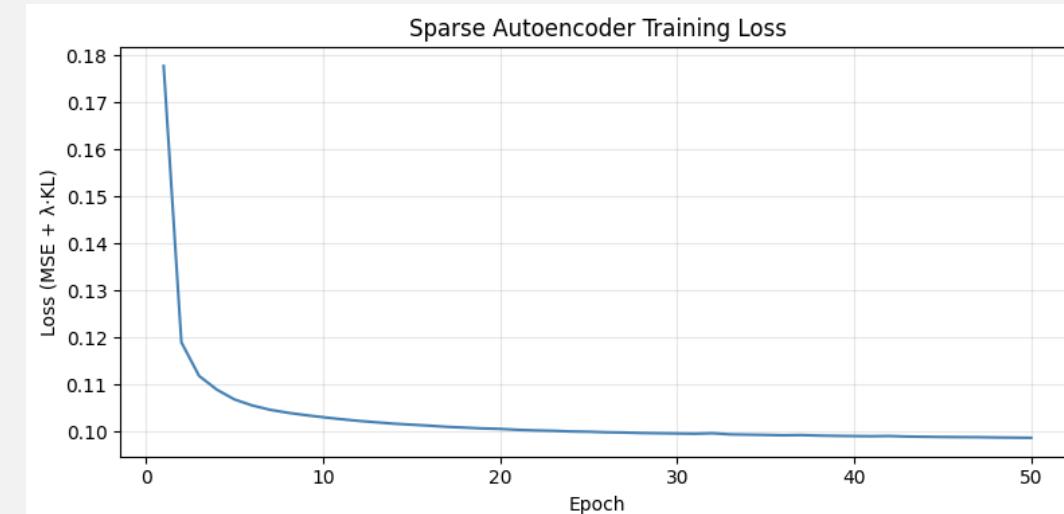
The ROBUST model was EASIER to fool (5/5 vs 0/5).

The lazy model relies on color shortcuts – pixel noise targets shape, which it ignores.

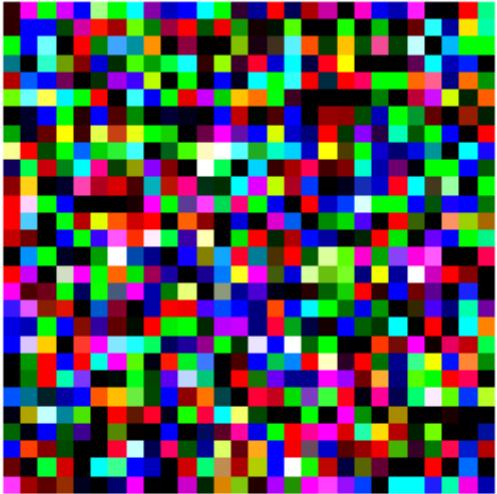
YES LAZY IS CRAZY IN LEARNING COLOURS BUT ROBUST IS ALSO CRAZZY IN GETTING ATTACKED

Task 6: The Decomposition: didn't use latent subspace index but yeah made it bit :

SAE: 2352 → 128 → 2352: WITH @ KLDIVERGENCE



Initial Random Noise
(shared for all neurons)

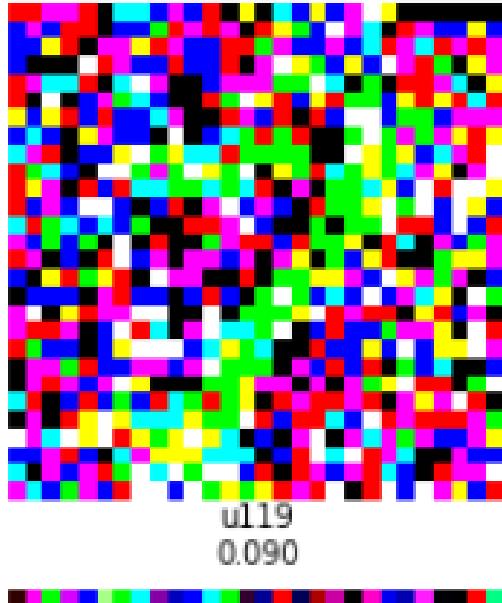


NOW HOOKED WRSPT TO BOTTLE NECK NEURONS

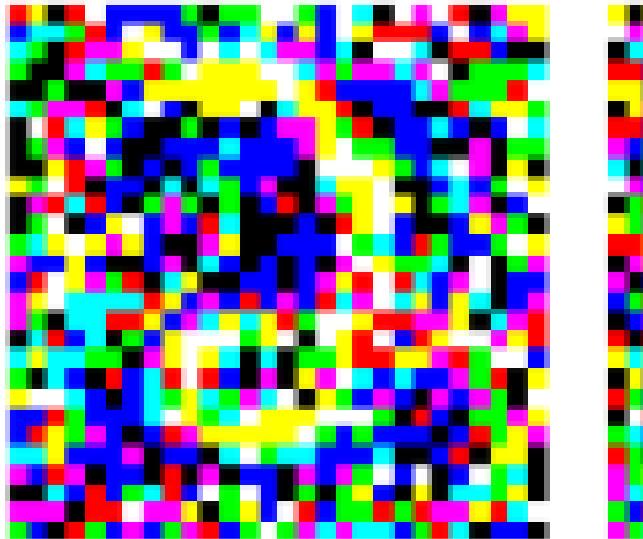
YES SOMETHING TAKES TIME: FROM COLORS, TEXTURES .. TO
POLYSEMATICITY WITH DIGITS @LSI



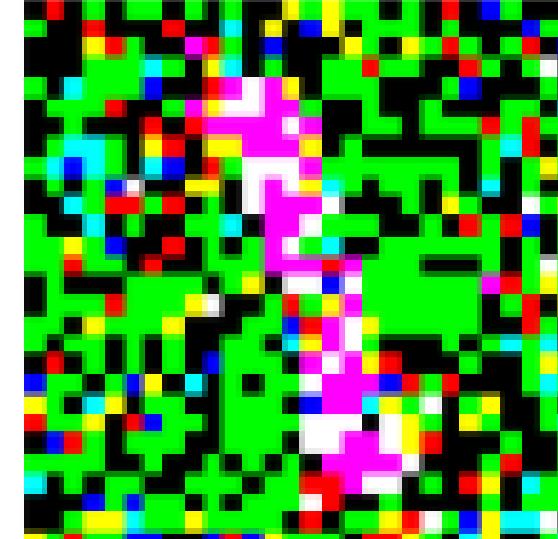
0.088



0.090



0.090



u126





THANK YOU SO MUCH TEAM PRECOG!! Task was
INSANE!!!!

my eyes are waiting to work under you team PRECOG
#pk sir #saralai ...IT's NOT OVER until I win!...

references: Selvaraj gradcam, ANDREJ KARPARTHY
convetjs,nptel lectures....